


# Autoencoder

14 Mar 2019

# Definition

Autoencoders are neural networks that are trained to copy their inputs to their outputs.

- Usually constrained in particular ways to make this task more difficult.



# Encoder and Decoder


- Encoder: Encodes its input  $x_i$  into a hidden representation  $h$

$$h = g(W_1 x_i + b)$$

- Decodes the input again from this hidden representation

$$\hat{x}_i = f(W_2 h + c)$$

- The model is trained to minimize a loss function which will ensure that  $\hat{x}_i$  is close to  $x_i$




# Undercomplete autoencoder

## 1. $\dim(h) < \dim(x_i)$

- Network must model  $x$  in lower dim. space + map latent space accurately back to input space.

Encoder network:

- hidden layer “compresses” the input
- will compress well only for the training distribution




$$h = g(W_1 x_i + b)$$

$$\hat{x}_i = f(W_2 h + c)$$

# Undercomplete autoencoder

1.  $\dim(h) < \dim(x_i)$

If network has only linear transformations, encoder learns PCA. With typical non-linearities, network learns generalized, more powerful version of PCA.




$$h = g(W_1 x_i + b)$$

$$\hat{x}_i = f(W_2 h + c)$$

# Link between PCA and Autoencoders


- the encoder part of an autoencoder is equivalent to PCA if we
  - use a linear encoder
  - use a linear decoder
  - use squared error loss function
  - normalize the inputs to

$$\hat{x}_{ij} = \frac{1}{\sqrt{m}} \left( x_{ij} - \frac{1}{m} \sum_{k=1}^m x_{kj} \right)$$



# Overcomplete Autoencoder

2.  $\dim(h) > \dim(x_i)$
- no compression in hidden layer
  - each hidden unit could copy a different input component
- No guarantee that the hidden units will extract meaningful structure



$$h = g(W_1 x_i + b)$$


$$\hat{x}_i = f(W_2 h + c)$$

# Binary Inputs

$$x_{ij} \in \{0,1\}$$

Decoder: use logistic function

Encoder: use sigmoid function



$$h = g(W_1 x_i + b)$$

$$\hat{x}_i = f(W_2 h + c)$$




# Real Inputs

$$x_{ij} \in R$$

Decoder: use linear function

Encoder: use sigmoid function




$$h = g(W_1 x_i + b)$$

$$\hat{x}_i = f(W_2 h + c)$$

# Loss Function: Real Inputs

$$x_{ij} \in R$$

$$\min_{W_1, W_2, b, c} \frac{1}{m} \sum_{i=1}^m \sum_{j=1}^n (\hat{x}_{ij} - x_{ij})^2$$



$$h = g(W_1 x_i + b)$$


$$\hat{x}_i = f(W_2 h + c)$$

# Loss Function: Binary Inputs

$$x_{ij} \in \{0,1\}$$

- For a single n-dimensional ith input we can use the following loss function

$$\min \left\{ - \sum_{j=1}^n (x_{ij} \log \hat{x}_{ij} + (1 - x_{ij}) \log(1 - \hat{x}_{ij})) \right\}$$



$$h = g(W_1 x_i + b)$$

$$\hat{x}_i = f(W_2 h + c)$$

# Regularization in autoencoders

- While poor generalization could happen even in undercomplete autoencoders it is an even more serious problem for overcomplete auto encoders
- To avoid poor generalization, we need to introduce regularization

# Denoising Autoencoders

Corrupts the input data using a


1. probabilistic process  
( $P(\tilde{x}_{ij}|x_{ij})$ ) before feeding  
it to the network

$$P(\tilde{x}_{ij} = 0|x_{ij}) = q$$

$$P(\tilde{x}_{ij} = x_{ij}|x_{ij}) = 1 - q$$


2. Gaussian additive noise

$$\tilde{x}_{ij} = x_{ij} + \mathcal{N}(0,1)$$




# Denoising Autoencoders

- the objective is to reconstruct the original (uncorrupted)  $x_i$
- By having to remove noise, model must know difference between noise and actual image.



# Denoising Autoencoder

- The corrupting function  $\mathcal{C}(\cdot)$  can corrupt in any direction
- autoencoder must learn "location" of data manifold and its distribution  $p_{data}(x)$



# Application of AE

- Task: Hand-written digit recognition



# Contractive Autoencoders

- Contractive Autoencoders are explicitly encouraged to learn a manifold through their loss function.
- Desirable property: Points close to each other in input space maintain that property in the latent space.
- objective is to have a robust learned representation which is less sensitive to small variation in the data.
- We wish to extract features that only reflect variations observed in the training set -- we'd like to be invariant to the other variations
- Robustness of the representation for the data is done by **applying a penalty term to the loss function. The penalty term is Frobenius norm of the Jacobian matrix.**

# Contractive Autoencoders

- Desirable property: Points close to each other in input space maintain that property in the latent space.
- This will be true if  $f(x) = h$  is continuous, has small derivatives.
- Robustness of the representation for the data is done by **applying a penalty term to the loss function.** The penalty term is Frobenius norm of the Jacobian matrix.

$$L(\theta) + \Omega(\theta)$$

$$\Omega(\theta) = \|J_x(h)\|_F^2$$

$J_x(h)$  is a Jacobian of the encoder.

# Jacobian and Frobenius Norm

- If the input has  $n$  dimensions and the hidden layer has  $k$  dimensions then

$$J_x(h) = \begin{bmatrix} \frac{\partial h_1}{\partial x_1} & \dots & \frac{\partial h_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial h_k}{\partial x_1} & \dots & \frac{\partial h_k}{\partial x_n} \end{bmatrix}$$

- The Frobenius Norm for a matrix  $M$ :

$$\|J_x(h)\|_F^2 = \sum_{j=1}^n \sum_{l=1}^k \frac{\partial h_l}{\partial x_j}$$

- Consider  $\frac{\partial h_1}{\partial x_1}$
- $\frac{\partial h_1}{\partial x_1} = 0$  means that this neuron is not very sensitive to variations cutting-edge the input  $x_1$ .
- $L(\theta)$  capture important variations in data
- $\Omega(\theta)$  do not capture variations in data

- Called contractive because they contract neighborhood of input space into smaller, localized group in latent space.
- This contractive effect is designed to only occur locally.
- The Jacobian Matrix will see most of its eigenvalues drop below 1  $\rightarrow$  contracted directions
- But some directions will have eigenvalues (significantly) above 1  $\rightarrow$  directions that explain most of the variance in data

# The Big Idea of Regularized Autoencoders

- Previous slides underscore the central balance of regularized autoencoders:
- Be sensitive to inputs (reconstruction loss) → generate good reconstructions of data drawn from data distribution
- Be insensitive to inputs (regularization penalty) → learn actual data distribution

# Sparse Autoencoders

- A hidden neuron with sigmoid activation will have values between 0 and 1
- We say that the neuron is activated when its output is close to 1 and not activated when its output is close to 0.
- A sparse autoencoder tries to ensure the neuron is inactive most of the times.

The average value of the activation of a neuron  $l$  is given by

$$\hat{\rho}_l = \frac{1}{m} \sum_{i=1}^m h(x_i)_l$$

If the neuron  $l$  is sparse (i.e. mostly inactive) then


$$\hat{\rho}_l \rightarrow 0$$

A sparse autoencoder uses a small value of sparsity parameter  $\rho$  (say, 0.005) and tries to enforce the constraint  $\hat{\rho}_l = \rho$

- One way of ensuring this is to add the following term to the objective function

$$\Omega(\theta) = \sum_{l=1}^k \rho \log \frac{\rho}{\hat{\rho}_l} + (1 - \rho) \log \frac{1 - \rho}{1 - \hat{\rho}_l}$$

# The sparsity function



The function will reach its minimum value(s) when  $\hat{\rho}_l = \rho$ .



# Sparse Autoencoders

- A sparse autoencoder involves a sparsity penalty  $\Omega(h)$  on the code layer  $h$ , in addition to the reconstruction error:

$$L(x, g(f(x))) + \Omega(h)$$

- Regularized maximum likelihood corresponds to maximizing  $p(\theta | x)$ , which is equivalent to maximizing

$$\log p(x|\theta) + \log p(\theta)$$

- The  $\log p(x | \theta)$  term is the usual data log-likelihood term and the  $\log p(\theta)$  term, the log-prior over parameters, incorporates the preference over particular values of  $\theta$ .

$$\Omega(\theta) = \sum_{l=1}^k \rho \log \frac{\rho}{\hat{\rho}_l} + (1 - \rho) \log \frac{1 - \rho}{1 - \hat{\rho}_l}$$

- Can be re-written as

$$\begin{aligned} & \Omega(\theta) \\ &= \sum_{l=1}^k \rho \log \rho - \rho \log \hat{\rho}_l + (1 - \rho) \log(1 - \rho) - (1 - \rho) \log(1 - \hat{\rho}_l) \end{aligned}$$

$$\Omega(\theta) = \sum_{l=1}^k \rho \log \frac{\rho}{\hat{\rho}_l} + (1 - \rho) \log \frac{1 - \rho}{1 - \hat{\rho}_l}$$

Can be re-written as

$$\Omega(\theta) = \sum_{l=1}^k \rho \log \rho - \rho \log \hat{\rho}_l + (1 - \rho) \log(1 - \rho) - (1 - \rho) \log(1 - \hat{\rho}_l)$$

By Chain rule:

$$\frac{\partial \Omega(\theta)}{\partial W} = \frac{\partial \Omega(\theta)}{\partial \hat{\rho}} \cdot \frac{\partial \hat{\rho}}{\partial W}$$

$$\frac{\partial \Omega(\theta)}{\partial \hat{\rho}} = \left[ \frac{\partial \Omega(\theta)}{\partial \hat{\rho}_1}, \frac{\partial \Omega(\theta)}{\partial \hat{\rho}_2}, \dots, \frac{\partial \Omega(\theta)}{\partial \hat{\rho}_k} \right]^T$$

For each neuron  $l \in 1 \dots k$  in hidden layer, we have

$$\frac{\partial \Omega(\theta)}{\partial \hat{\rho}_l} = -\frac{\rho}{\hat{\rho}_l} + \frac{(1 - \rho)}{1 - \hat{\rho}_l}$$

and  $\frac{\partial \hat{\rho}_l}{\partial W} = \mathbf{x}_i (g'(W^T \mathbf{x}_i + \mathbf{b}))^T$  (see next slide)

- Now,

$$\hat{\mathcal{L}}(\theta) = \mathcal{L}(\theta) + \Omega(\theta)$$

- $\mathcal{L}(\theta)$  is the squared error loss or cross entropy loss and  $\Omega(\theta)$  is the sparsity constraint.
- We already know how to calculate  $\frac{\partial \mathcal{L}(\theta)}{\partial W}$
- Let us see how to calculate  $\frac{\partial \Omega(\theta)}{\partial W}$ .
- Finally,

$$\frac{\partial \hat{\mathcal{L}}(\theta)}{\partial W} = \frac{\partial \mathcal{L}(\theta)}{\partial W} + \frac{\partial \Omega(\theta)}{\partial W}$$

(and we know how to calculate both terms on R.H.S)

## Derivation

$$\frac{\partial \hat{\rho}}{\partial W} = \begin{bmatrix} \frac{\partial \hat{\rho}_1}{\partial W} & \frac{\partial \hat{\rho}_2}{\partial W} & \cdots & \frac{\partial \hat{\rho}_k}{\partial W} \end{bmatrix}$$

For each element in the above equation we can calculate  $\frac{\partial \hat{\rho}_l}{\partial W}$  (which is the partial derivative of a scalar w.r.t. a matrix = matrix). For a single element of a matrix  $W_{jl}$ :-

$$\begin{aligned} \frac{\partial \hat{\rho}_l}{\partial W_{jl}} &= \frac{\partial \left[ \frac{1}{m} \sum_{i=1}^m g(W_{:,l}^T \mathbf{x}_i + b_l) \right]}{\partial W_{jl}} \\ &= \frac{1}{m} \sum_{i=1}^m \frac{\partial \left[ g(W_{:,l}^T \mathbf{x}_i + b_l) \right]}{\partial W_{jl}} \\ &= \frac{1}{m} \sum_{i=1}^m g'(W_{:,l}^T \mathbf{x}_i + b_l) x_{ij} \end{aligned}$$

So in matrix notation we can write it as :

$$\frac{\partial \hat{\rho}_l}{\partial W} = \mathbf{x}_i (g'(W^T \mathbf{x}_i + \mathbf{b}))^T$$

# Representational Power, Layer Size and Depth

- Deeper autoencoders tend to generalize better and train more efficiently than shallow ones.
  - Common strategy: greedily pre-train layers and stack them
  - For contractive autoencoders, calculating Jacobian for deep networks is expensive. Good idea to do layer-by-layer.

# Applications of Autoencoders

- Dimensionality Reduction: Make high-quality, low-dimension representation of data
- Information Retrieval: Locate value in database which is just autoencoded key.